



# DEPARTMENT OF COMPUTER SCIENCE & ENGINEERING

Discover. Learn. Empower.

## Experiment 5

**Student Name:** Sambhav Mahajan

**UID:** 23BCS11290

**Branch:** B.E. C.S.E.

**Section/Group:** KRG-2B

**Semester:** 5th

**Date of Performance:**

**Subject Name:** ADBMS

**Subject Code:** 23CSP-333

**1. Aim:** To demonstrate the performance difference between a normal (logical) view and a materialized view in PostgreSQL by (1) generating a very large transaction dataset using generate\_series() and random(), (2) building both a normal view and a materialized view that compute total\_quantity\_sold, total\_sales, and total\_orders, and (3) measuring & comparing query execution times.

## 2. Objective:

1. Create a large transaction\_data table populated using generate\_series() and random().
2. Create a normal view sales\_summary and a materialized view sales\_summary\_mat that aggregate total\_quantity\_sold, total\_sales, and total\_orders.
3. Measure and compare execution time and resource usage for querying each view, and for refreshing the materialized view.
4. Learn when to use materialized views vs normal views in OLAP/analytics scenarios.
5. Assumption used: The script below generates 1,000,000 rows for id = 1 and 1,000,000 rows for id = 2 (total 2,000,000 rows).

## 3. DBMS Script:

```
-- =====  
-- Variant A: Minimal table (follows user's "id, value" schema)  
-- =====  
  
-- Drop table if exists  
DROP TABLE IF EXISTS transaction_data;  
  
-- Create table  
CREATE TABLE transaction_data (  
    id integer NOT NULL,  
    value numeric(12,2) NOT NULL -- sale amount for that row
```



# DEPARTMENT OF COMPUTER SCIENCE & ENGINEERING

Discover. Learn. Empower.

);

```
-- Insert 1,000,000 rows for id=1
INSERT INTO transaction_data (id, value)
SELECT 1, round((random() * 100)::numeric, 2)
FROM generate_series(1,1000000);
```

```
-- Insert 1,000,000 rows for id=2
INSERT INTO transaction_data (id, value)
SELECT 2, round((random() * 100)::numeric, 2)
FROM generate_series(1,1000000);
```

```
-- Update planner statistics
VACUUM ANALYZE transaction_data;
```

```
-- -----
-- Normal View (recomputes on each query)
-- -----
```

```
DROP VIEW IF EXISTS sales_summary;
```

```
CREATE VIEW sales_summary AS
SELECT
    id,
    COUNT(*) AS total_quantity_sold, -- assuming 1 item per row
    SUM(value) AS total_sales,
    COUNT(*) AS total_orders        -- one order per row in this minimal schema
FROM transaction_data
GROUP BY id;
```

```
-- -----
-- Materialized View (stores computed result)
-- -----
```

```
DROP MATERIALIZED VIEW IF EXISTS sales_summary_mat;
```

```
CREATE MATERIALIZED VIEW sales_summary_mat AS
SELECT
    id,
    COUNT(*) AS total_quantity_sold,
    SUM(value) AS total_sales,
    COUNT(*) AS total_orders
```



# DEPARTMENT OF COMPUTER SCIENCE & ENGINEERING

Discover. Learn. Empower.

```
FROM transaction_data
GROUP BY id
WITH NO DATA; -- create empty first, then refresh

-- Populate the materialized view
REFRESH MATERIALIZED VIEW sales_summary_mat;

-- Optional: index for faster lookups
CREATE UNIQUE INDEX IF NOT EXISTS sales_summary_mat_id_idx
ON sales_summary_mat (id);

-- =====
-- Variant B: More realistic schema with quantity & price
-- =====

-- Drop table if exists
DROP TABLE IF EXISTS transaction_data_real;

-- Create table
CREATE TABLE transaction_data_real (
    id      integer NOT NULL,
    quantity integer NOT NULL,
    price   numeric(12,2) NOT NULL,
    created_at timestamp with time zone DEFAULT now()
);

-- Insert 1,000,000 rows for id=1
INSERT INTO transaction_data_real (id, quantity, price)
SELECT 1,
       (floor(random()*5)+1)::int AS quantity,
       round((1 + random()*199)::numeric, 2) AS price
FROM generate_series(1,1000000);

-- Insert 1,000,000 rows for id=2
INSERT INTO transaction_data_real (id, quantity, price)
SELECT 2,
       (floor(random()*5)+1)::int AS quantity,
       round((1 + random()*199)::numeric, 2) AS price
```



# DEPARTMENT OF COMPUTER SCIENCE & ENGINEERING

Discover. Learn. Empower.

```
FROM generate_series(1,1000000);

-- Update planner statistics
VACUUM ANALYZE transaction_data_real;

-----
-- Normal View (recomputes on each query)
-----

DROP VIEW IF EXISTS sales_summary_real;

CREATE VIEW sales_summary_real AS
SELECT
    id,
    SUM(quantity) AS total_quantity_sold,
    SUM(price * quantity) AS total_sales,
    COUNT(*) AS total_orders
FROM transaction_data_real
GROUP BY id;

-----
-- Materialized View (stores computed result)
-----

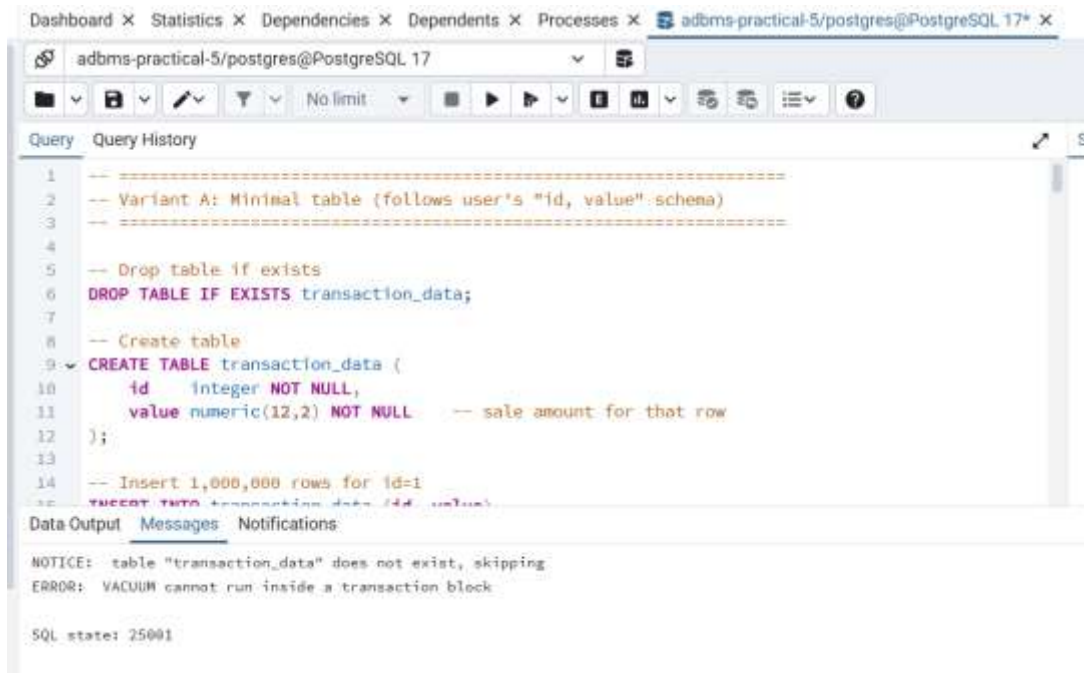
DROP MATERIALIZED VIEW IF EXISTS sales_summary_real_mat;

CREATE MATERIALIZED VIEW sales_summary_real_mat AS
SELECT
    id,
    SUM(quantity) AS total_quantity_sold,
    SUM(price * quantity) AS total_sales,
    COUNT(*) AS total_orders
FROM transaction_data_real
GROUP BY id
WITH NO DATA;

-- Populate the materialized view
REFRESH MATERIALIZED VIEW sales_summary_real_mat;

-- Optional: index for faster lookups
CREATE UNIQUE INDEX IF NOT EXISTS sales_summary_real_mat_id_idx
ON sales_summary_real_mat (id);
```

## 4. Output:



```

1  -- =====
2  -- Variant A: Minimal table (follows user's "id, value" schema)
3  -- =====
4
5  -- Drop table if exists
6  DROP TABLE IF EXISTS transaction_data;
7
8  -- Create table
9  CREATE TABLE transaction_data (
10     id integer NOT NULL,
11     value numeric(12,2) NOT NULL -- sale amount for that row
12 );
13
14 -- Insert 1,000,000 rows for id=1
15 INSERT INTO transaction_data (id, value)

```

Data Output Messages Notifications

NOTICE: table "transaction\_data" does not exist, skipping  
 ERROR: VACUUM cannot run inside a transaction block

SQL state: 25001

Figure 1

## 5. Learning Outcomes:

1. Use generate\_series() and random() to rapidly populate large test datasets in PostgreSQL.
2. Create and populate normal views and materialized views. Understand the syntactic difference: views are logical; materialized views store results.
3. Use EXPLAIN ANALYZE (and EXPLAIN (ANALYZE, BUFFERS)) and psql \timing to measure query execution time and to inspect query plans.
4. Compare trade-offs:
5. Normal view: always fresh, but expensive to compute repeatedly.
6. Materialized view: fast reads, but needs refresh (cost to recompute), and may be stale.
7. Apply indexing and VACUUM/ANALYZE to aid planner decisions & performance.
8. Decide when materialized views are appropriate (read-heavy analytics with acceptable staleness) vs when dynamic aggregation is required.